


Real-World Federation of Autonomous Kubernetes in an Interconnected Continuum

Giuseppe Zangari
Aruba Kube

Fulvio Risso
Politecnico di Torino


Can our clusters keep up?

Research is hitting a compute wall



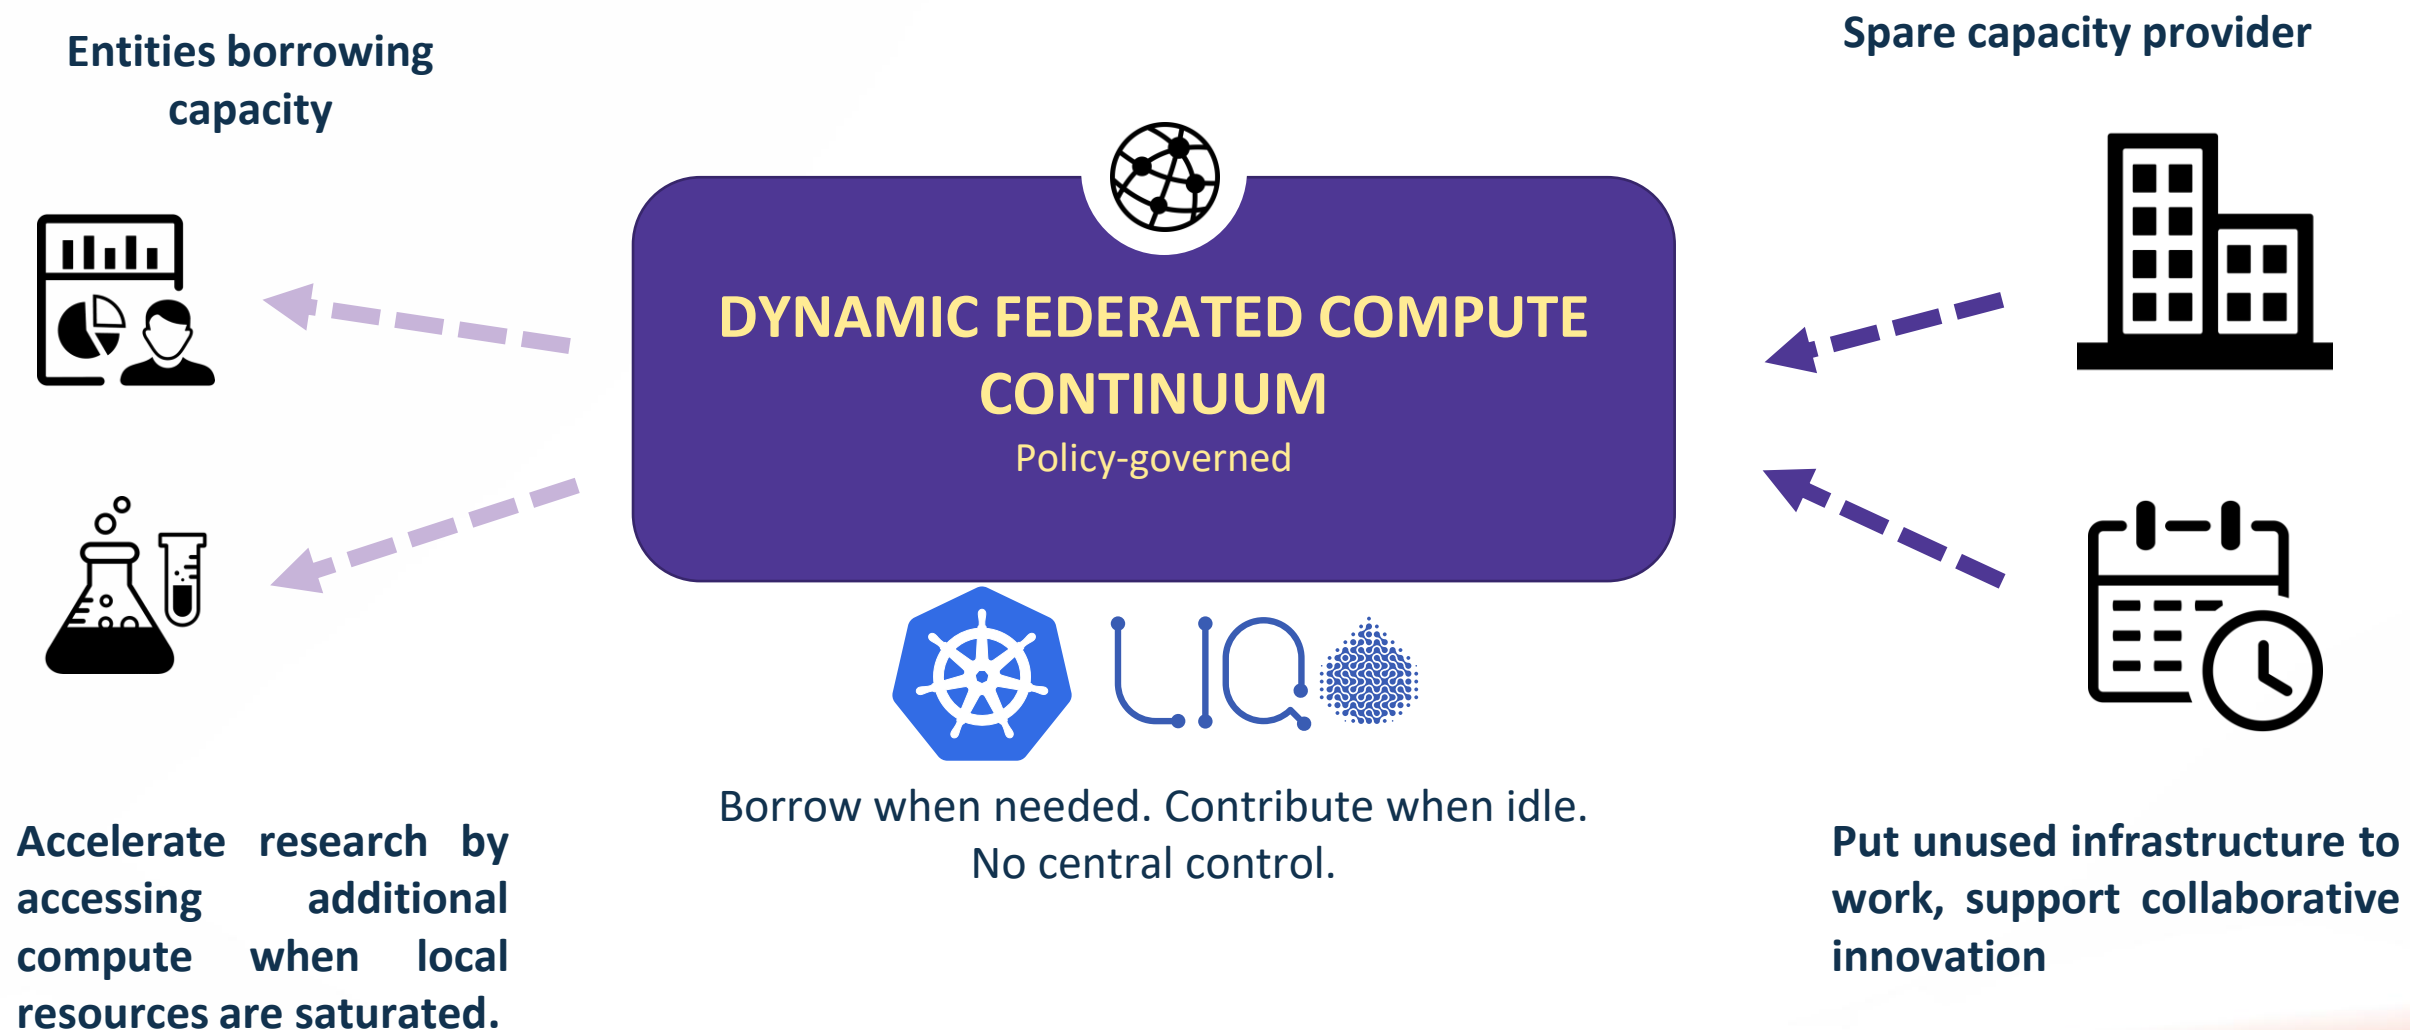
From AI to climate models, **modern science demands more CPU, GPU, and I/O**. Without flexible, scalable infrastructure, **compute becomes a bottleneck—and innovation slows down**.

The use of compute is discontinuous



Compute usage is inherently discontinuous, intense bursts during coding-heavy weeks, lulls when work shifts to design, meetings, or data prep. Provisioning for the “average” load leaves **systems swinging between congestion and wasted capacity**.

Scaling and Sharing beyond clusters walls



AGER Today: Live, Running, Federated

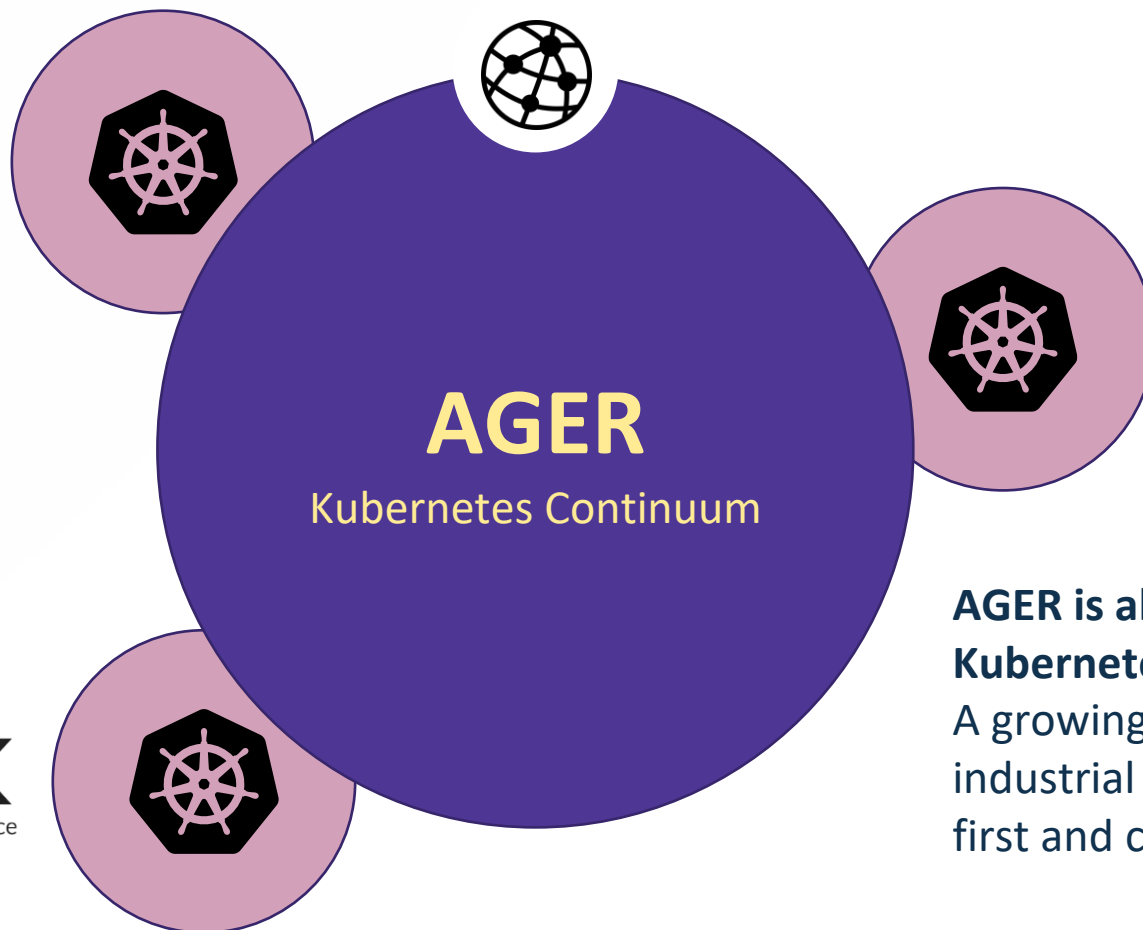


**Politecnico
di Torino**

Bare-metal Cluster
Turin



Bare-metal Cluster
Bologna



aruba.it



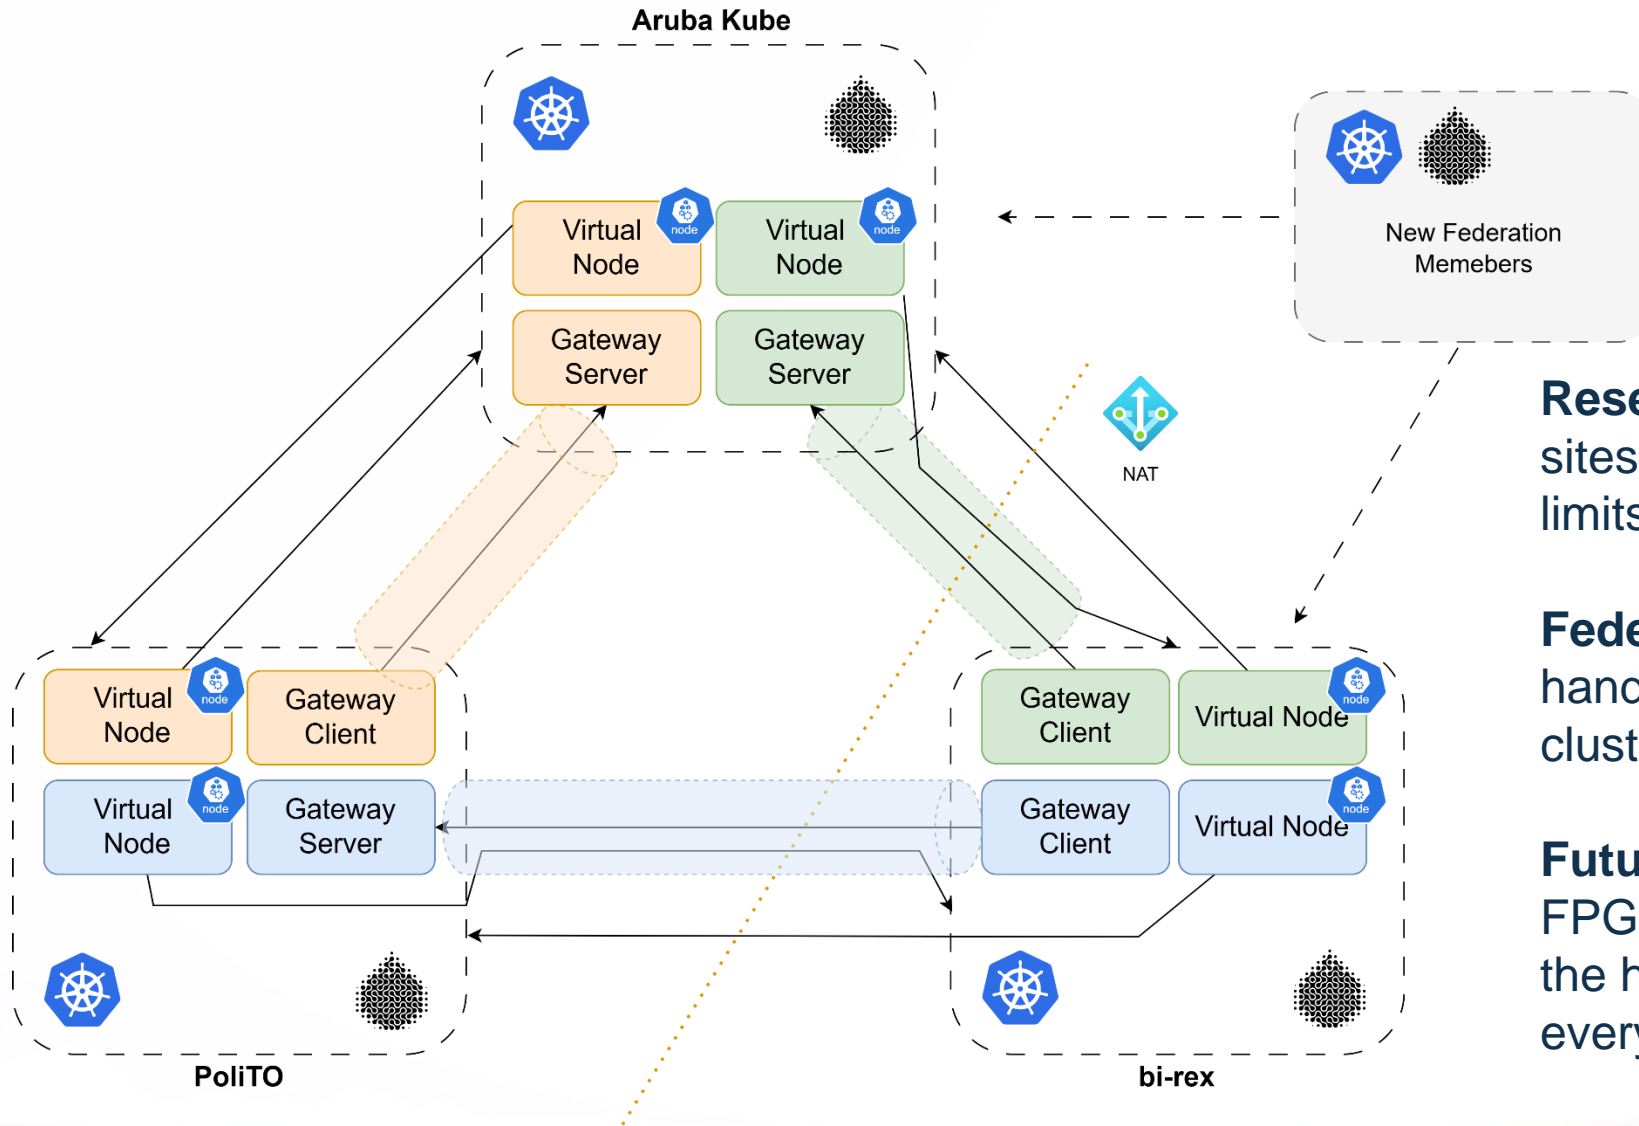
Aruba Managed Kubernetes
Bergamo

AGER is already live, connecting production-grade Kubernetes clusters across Italy.

A growing federation where academic and industrial partners scale and share compute, policy-first and cloud-agnostic.

We are actively seeking new contributors: bring your cluster, offer spare capacity, or burst when you need it.

What powers AGER? Transparent federation via Liko 1.0



Research without walls Run across sites as if local. No more queue-induced limits on collaboration or scale.

Federation-as-code One install, and Liko handles peering, policies, and multi-cluster identity, zero disruption.

Future-ready compute GPUs today, FPGAs tomorrow, Kubernetes abstracts the hardware, Liko makes it usable everywhere.

Run Federated. Feel Local. Stay in Control.

```
fra in 🌀ager-arubakube in ~
> liqctl info
Local installation info
Cluster ID:      ager-arubakube
Version:         v1.0.0
K8s API server: https://95.110.142.35:6443
Cluster labels
  liq.io/provider:      aruba
  topology.kubernetes.io/zone: bergamo
```

```
Installation health
✓ Ligo is healthy
```

```
Active peerings
bi-rex-cluster
  Role:      ConsumerAndProvider
  Networking status: Healthy
  Authentication status: Healthy
  Offloading status: Healthy
poli-cladispe
  Role:      ConsumerAndProvider
  Networking status: Healthy
  Authentication status: Healthy
  Offloading status: Healthy
```

Deploy as usual – Use your existing kubectl, Helm, or CI pipeline, no need to learn new tools.

Burst automatically – When local resources are full, workloads transparently overflow to trusted AGER peers.

Stay in control – You define what can run where, with quotas, labels, and namespaces enforced by policy. You decide what to share.

See what's happening – Unified observability lets you track jobs and resource usage across clusters.

Scale with us: From core Partners to a Community Fabric

AGER next steps

AGER Bootstrap

May 2025

Brainstorm & Process Definition with new partners

Q3 2025

Scale & Improve

Q4 2025 onwards

AGER is live, but not finished. You can help shape how this federation grows: technically, organizationally, and ethically.

We're building the policies together.

Quotas, access, incentives, fairness, partners will co-design the rules, not just follow them.


You don't need to commit today. Peer in, test workloads, explore governance. The door is open—no vendor, no lock-in.

Why now? Because the foundation is running—and it's solid. But what it becomes is up to the community that joins next.

Join not just to run workloads, join to invent the future of shared academic infrastructure.

Let's innovate together!

Fulvio Riso



Giuseppe
Zangari

