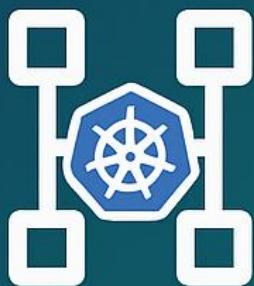




# Master Kubernetes Networking

From Basics to Advanced  
eBPF & Enterprise SDN



42 slides of pure  
networking mastery

CNI Plugins • Service Types • OVN-K8s  
eBPF • Production Tips

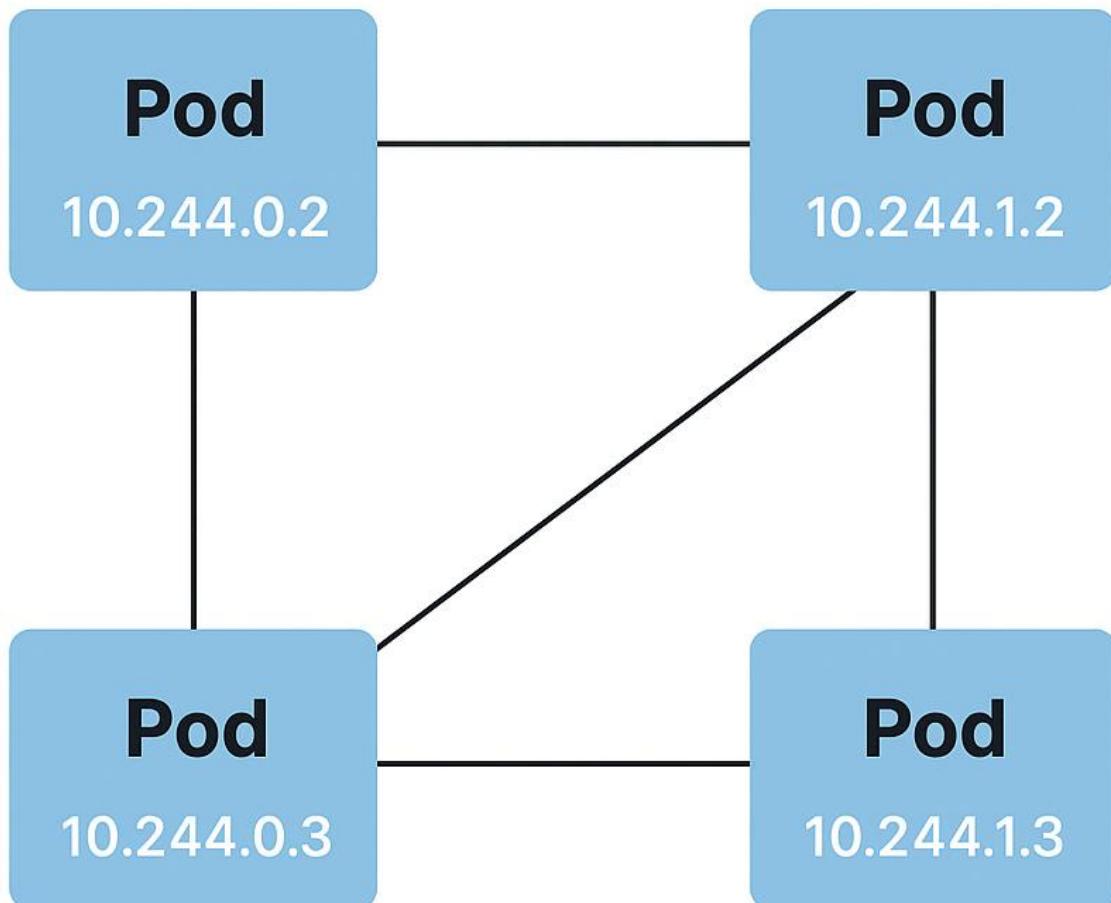
Swipe for complete networking guide →

# What's Inside This Guide:

- Kubernetes networking fundamentals
- CNI plugins deep dive
- eBPF implementation secrets
- Real-world troubleshooting
- Production best practices
- Advanced load balancing

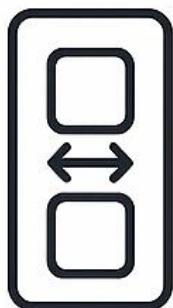
# The K8s Networking Foundation

- Flat address space for all pods
- Each pod gets unique IP address
- No NAT required between pods
- Containers share pod IP via localhost

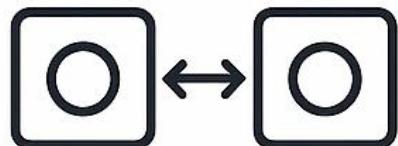


# THE FOUR PILLARS OF K8s NETWORKING

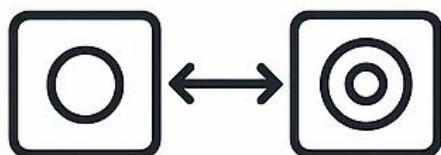
4 Core Communication Types:



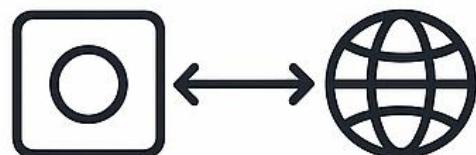
**Intra-pod**  
(container  $\leftrightarrow$  container)



**Pod-to-pod  
communication**



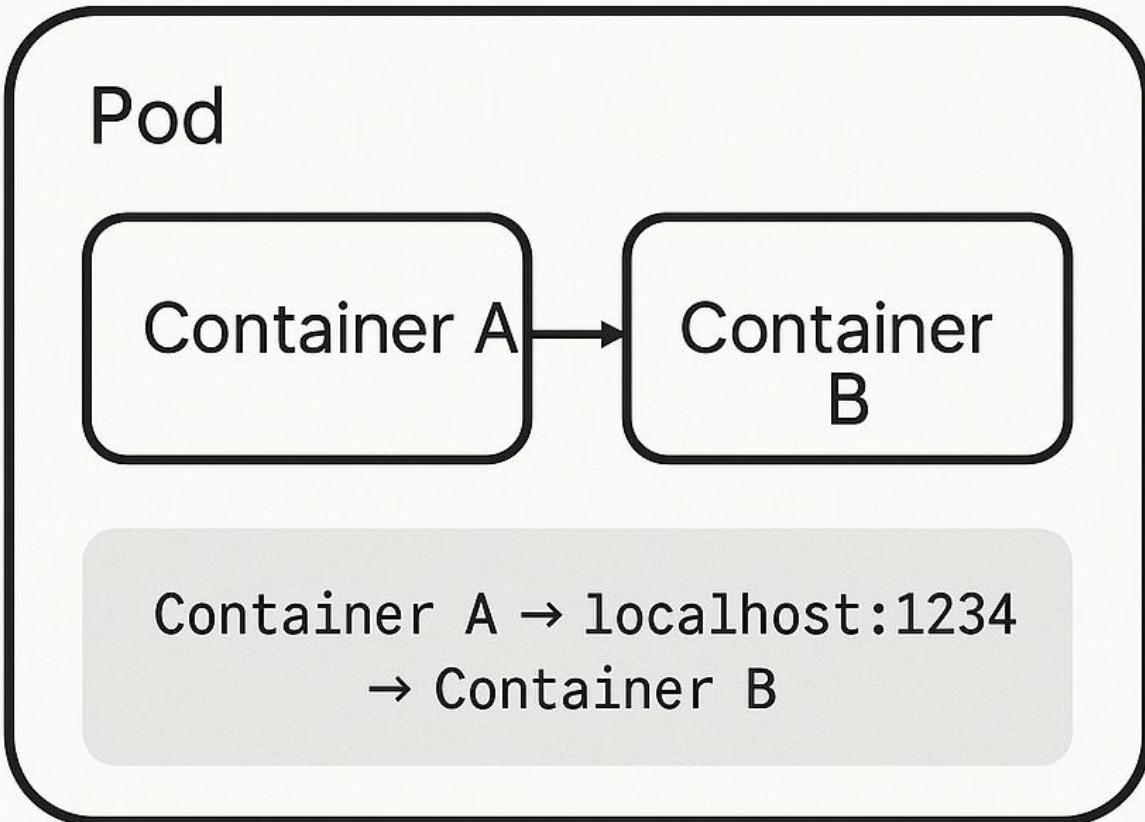
**Pod-to-service  
communication**



**External access  
patterns**

Master these = Master K8s networking

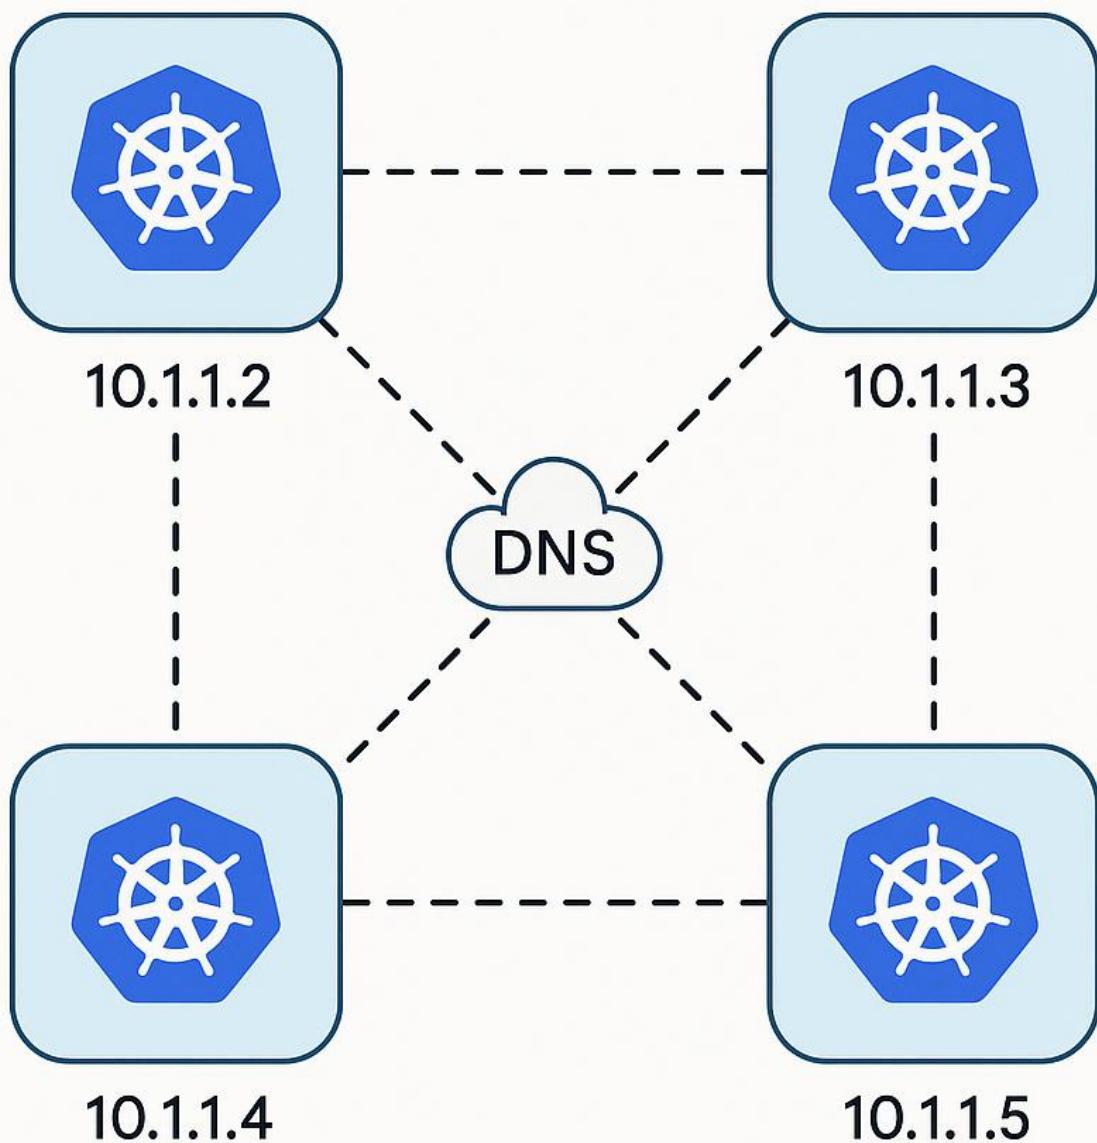
# Inside the Pod Network



Containers share  
network namespace

**Pro Tip:**  
No port conflicts between pods!

# Direct Pod Communication



- Network-visible IP addresses
- No NAT/tunnels/proxies needed
- Standard DNS works out of the box
- Well-known ports for easy config

**Seamless app migration**

# Kubernetes Services Explained

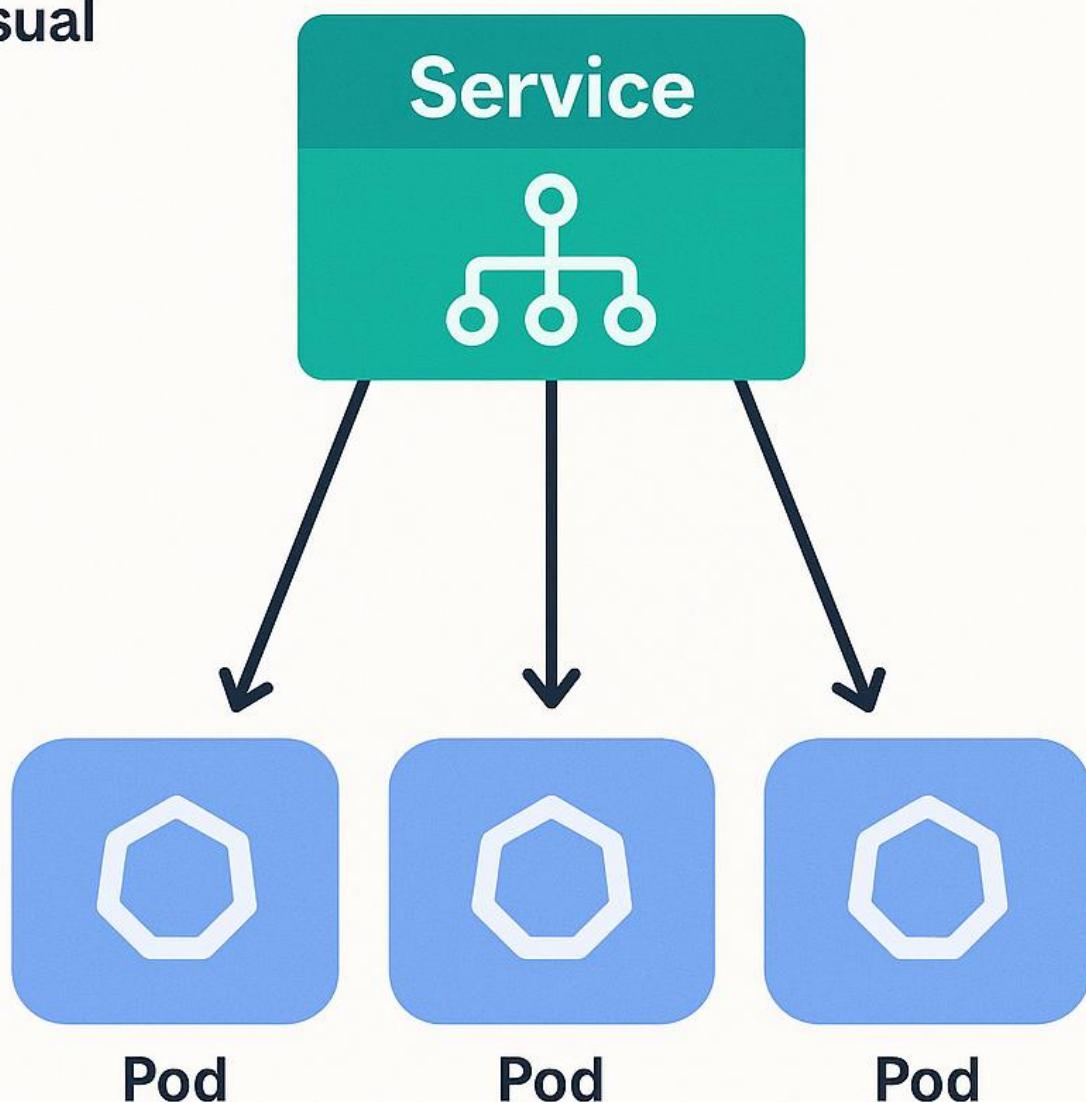
**Problem:** Pods are ephemeral, IPs change

**Solution:** Services provide stable endpoints

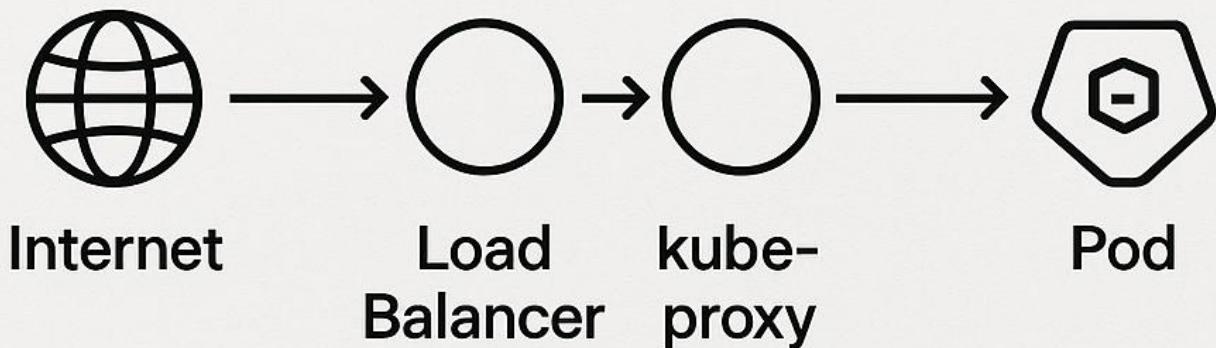
**Benefits**

- Automatic load balancing
- Service discovery
- Health checking

**Visual**



# Reaching Your Apps

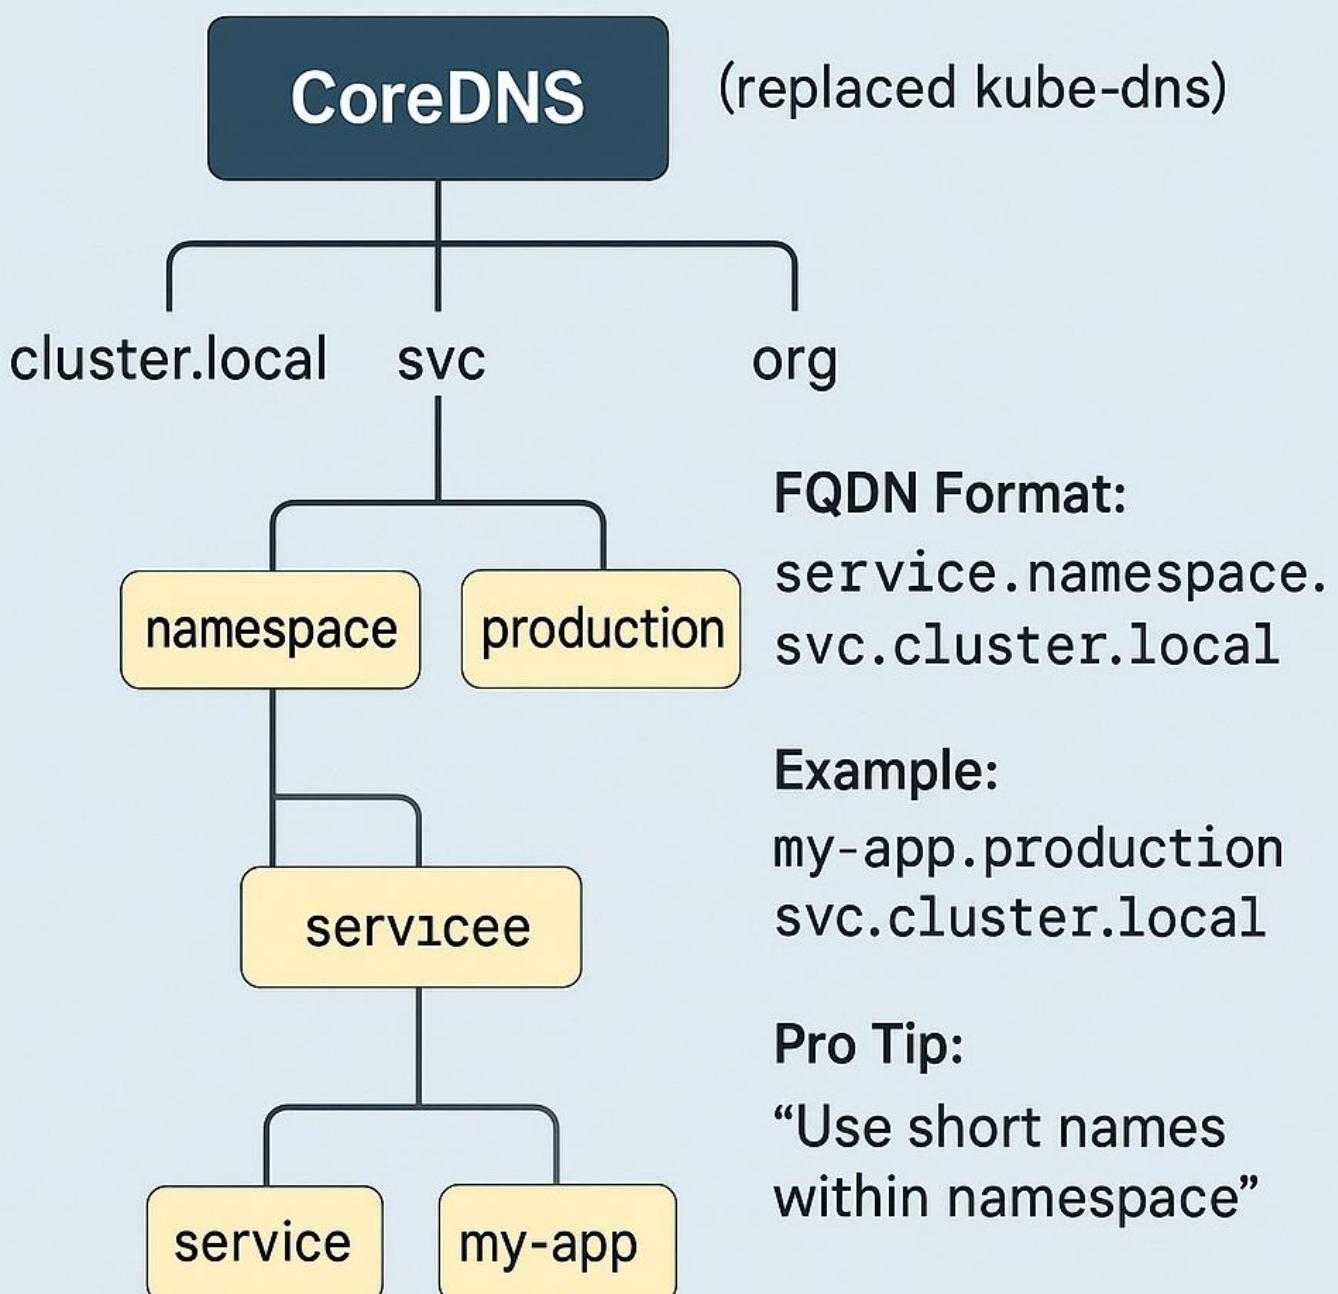


**Challenge:** Double-hop routing overhead

**Solutions:** NodePort,  
LoadBalancer, Ingress

**Performance tip:** Consider service mesh for optimization

# Name Resolution Made Simple



# DNS Setup in Action

- Headless service creation

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  clusterIP: None
  ports: 80
  targetPort: 9376
```

- Pod hostname configuration

```
apiVersion: v1
kind: Pod
metadata: name-my-pod
spec:
  hostname: busybox
  containers:
    - name: busybox
      image: busybox
```

- FQDN resolution

```
$ host my-pod.my-service.default.svc.
cluster.local has address 10.1.0.120
```

---

Troubleshooting: Check /etc/resolv.conf in pods

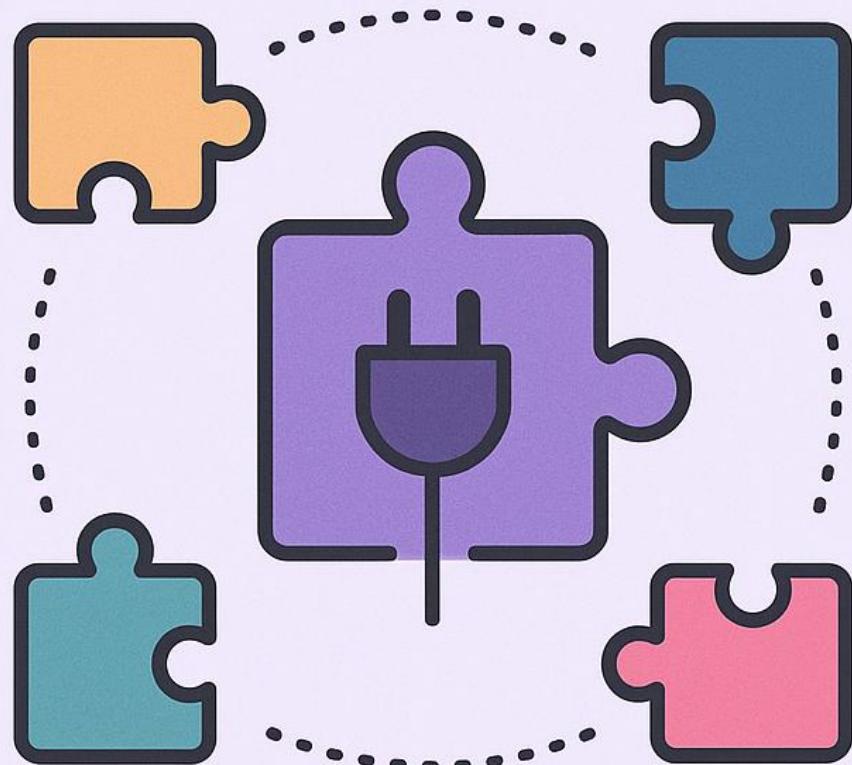
# The Plugin Revolution

## Why needed:

Diverse networking requirements

## Two main types:

- Kubenet (basic, legacy)
- CNI (modern, extensible)



**Choose what fits  
your needs**

# Built on Linux Networking

- Network namespaces
- Virtual Ethernet (veth) pairs
- Bridges and routing
- IP addresses and ports

Understanding Linux =  
Understanding K8s

# veth: The Virtual Cable

## Purpose

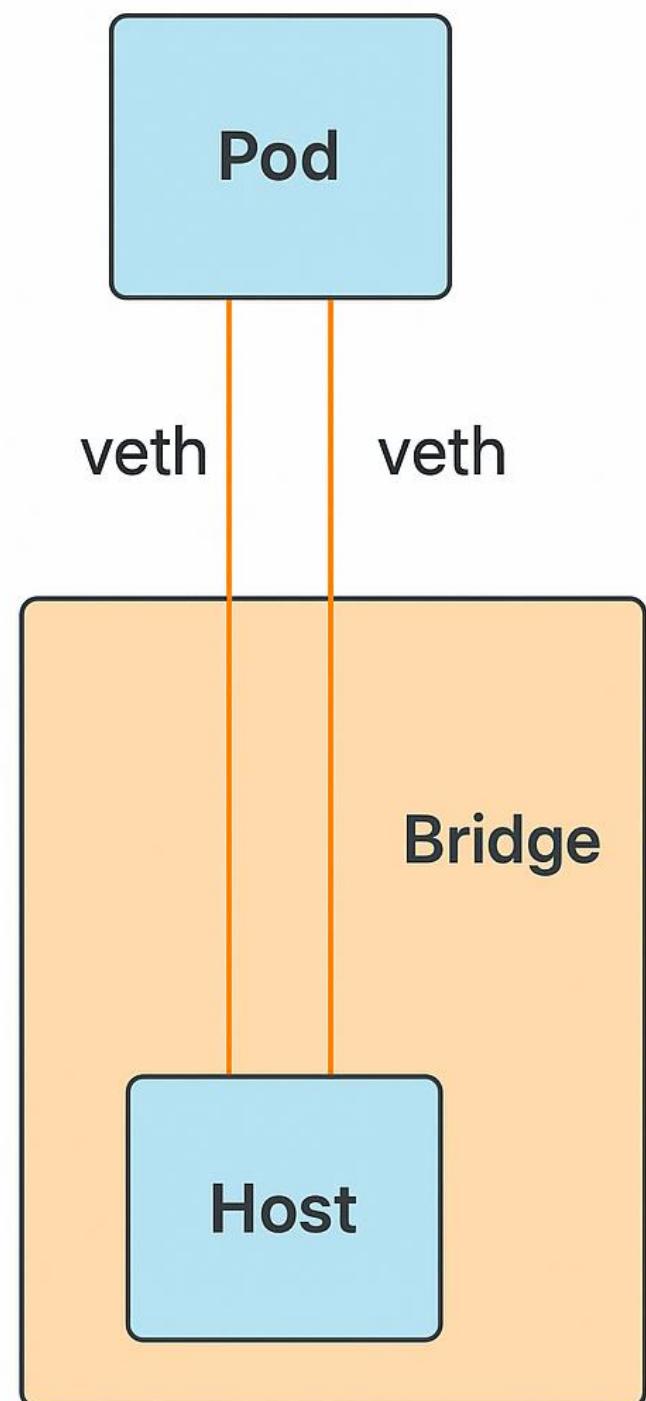
Connect containers to host network

## How it works

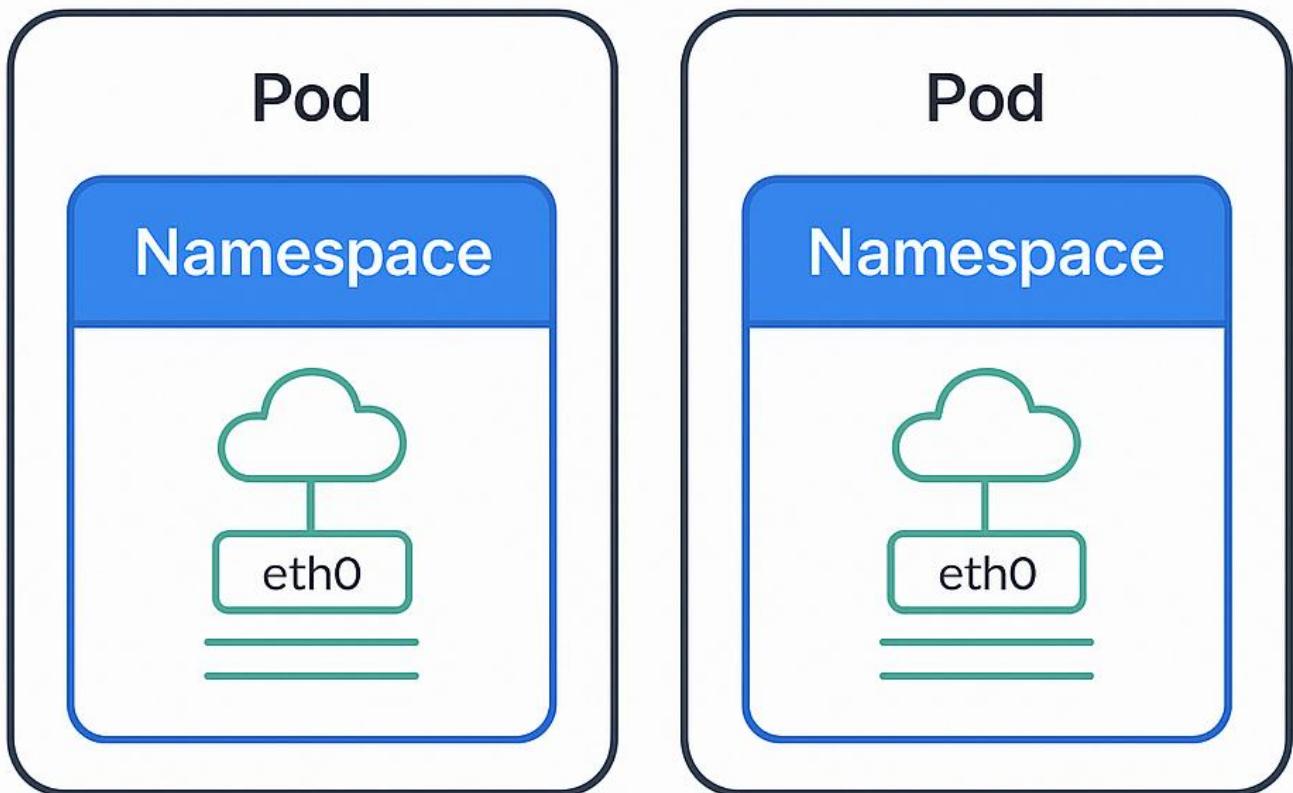
Paired virtual interfaces

## Use case

Pod  $\leftrightarrow$  Host communication



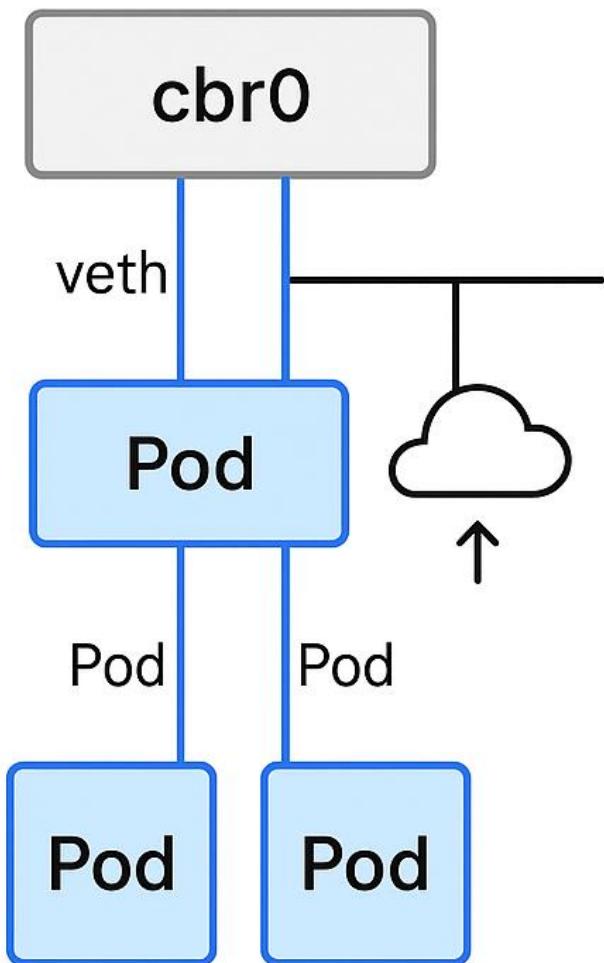
# Isolation Through Namespaces



- Network isolation
- No IP conflicts
- Security boundaries

Each pod = own namespace

# Kubenet: Keep It Simple



Creates cbr0 bridge  
+ veth pairs

## Use cases:

- Single-node environments
- Cloud provider routing
- Learning/development

## Limitation:

“Basic functionality only”

# CNI: Container Network Interface

## Why it matters

Industry standard

## Used by:

Kubernetes

OpenShift

Mesos

AWS EKS



Kubernetes OpenShift Mesos AWS EKS

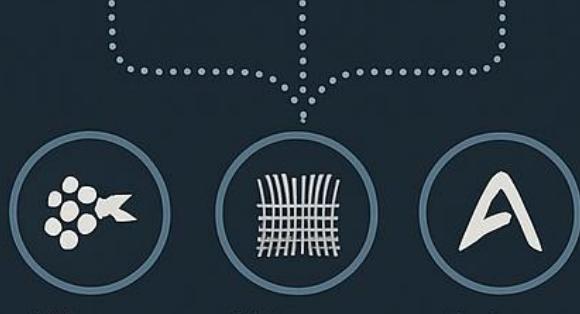
## Benefits:

- Vendor flexibility
- Rich plugin ecosystem
- Standardized interface

CNI



Calico Cilium Flannel



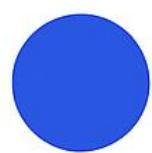
Weave Weave Antrea

# CNI Plugin Ecosystem



**Calico**

Layer 3 virtual  
networking



**Flannel**

Simple overlay  
networking



**Weave Net**

Zero-config  
networking

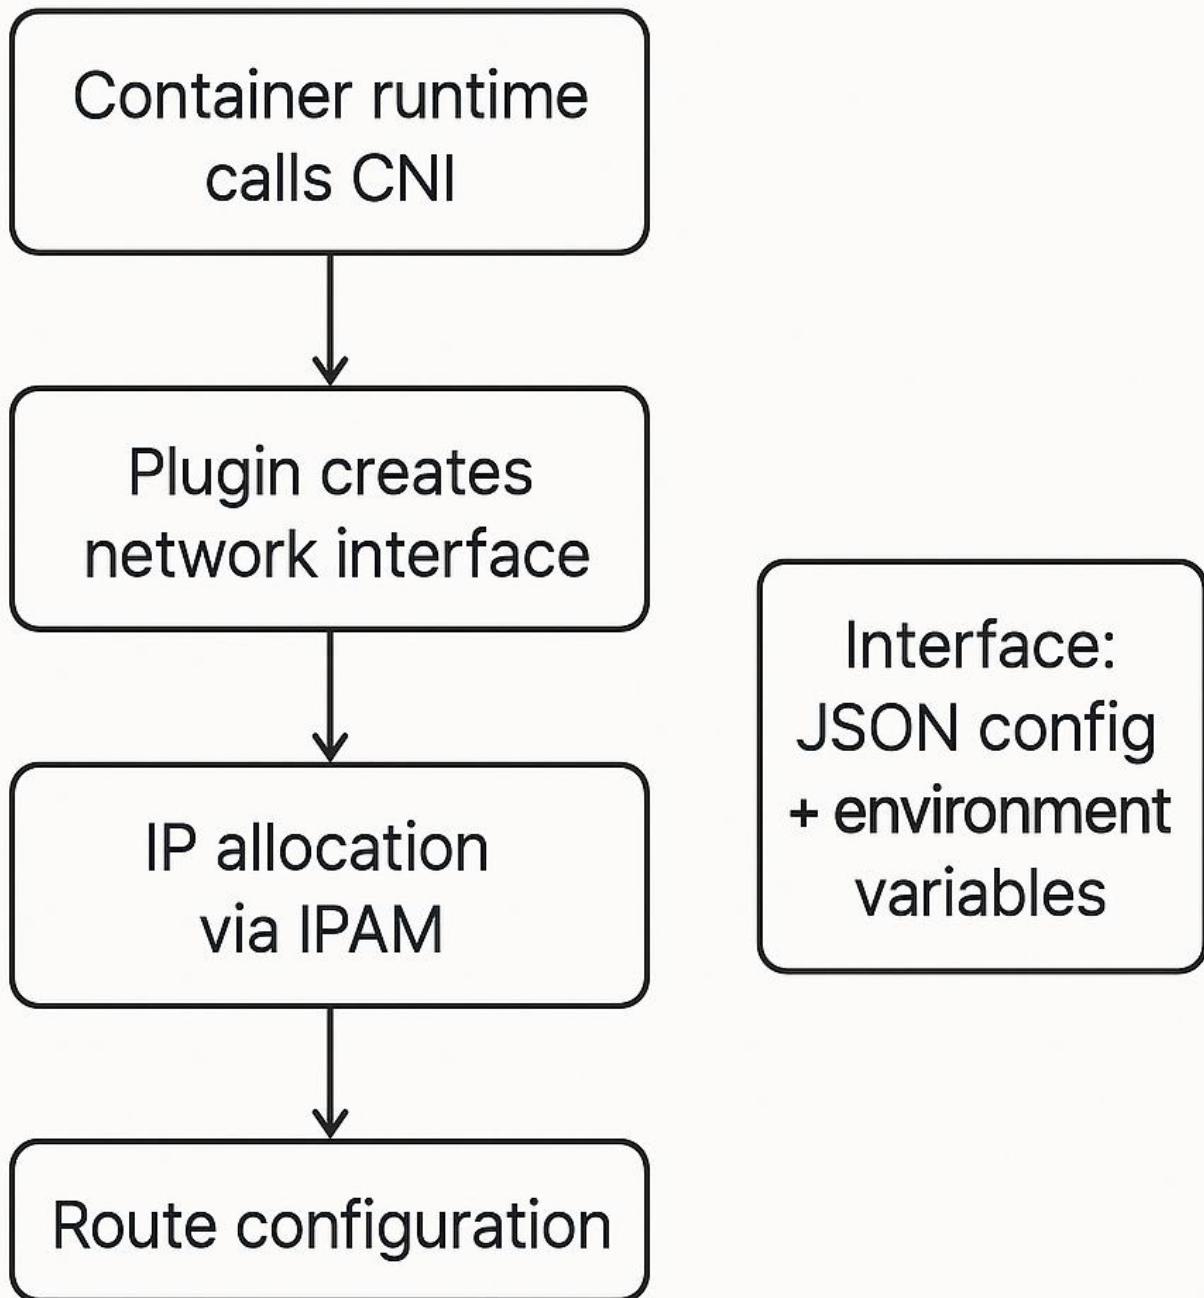


**Cilium**

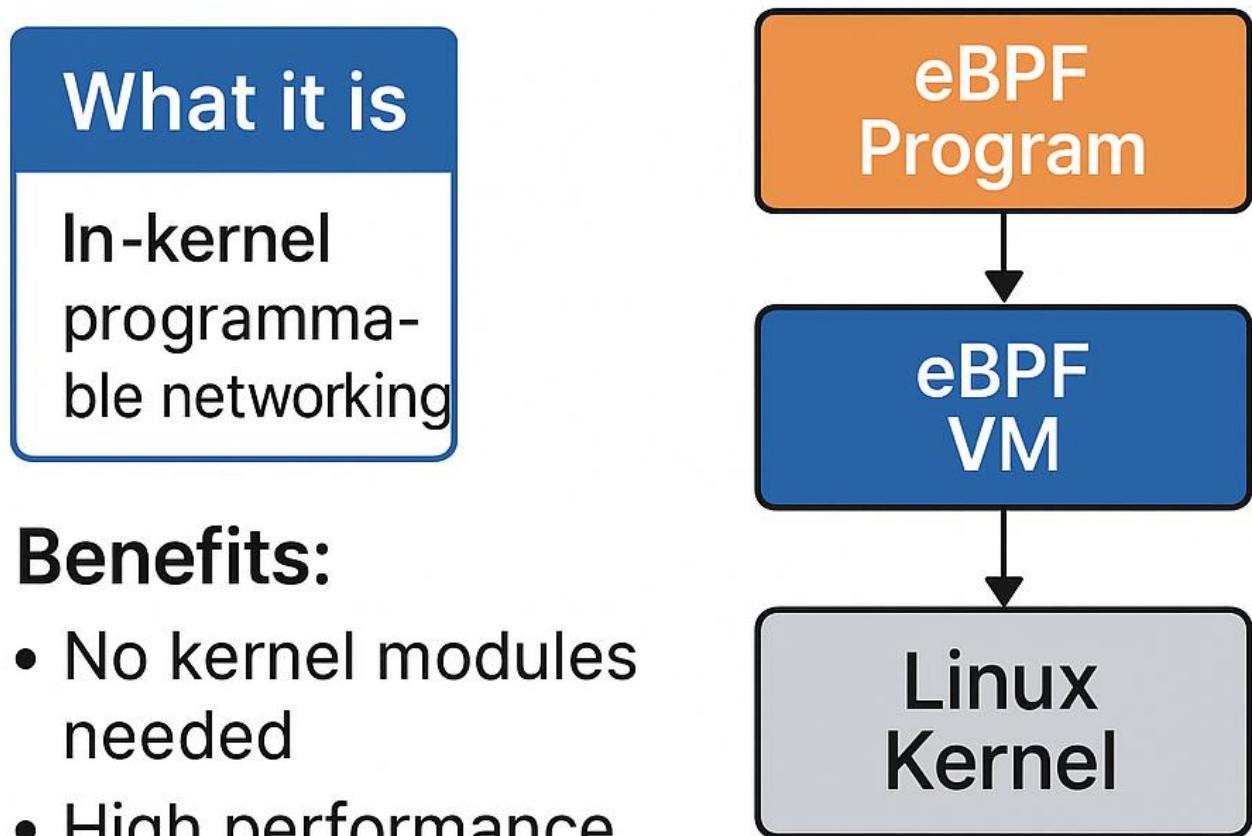
eBPF-powered  
networking

Each with specific strengths and use cases

# How CNI Plugins Work



# eBPF: The Linux Superpower

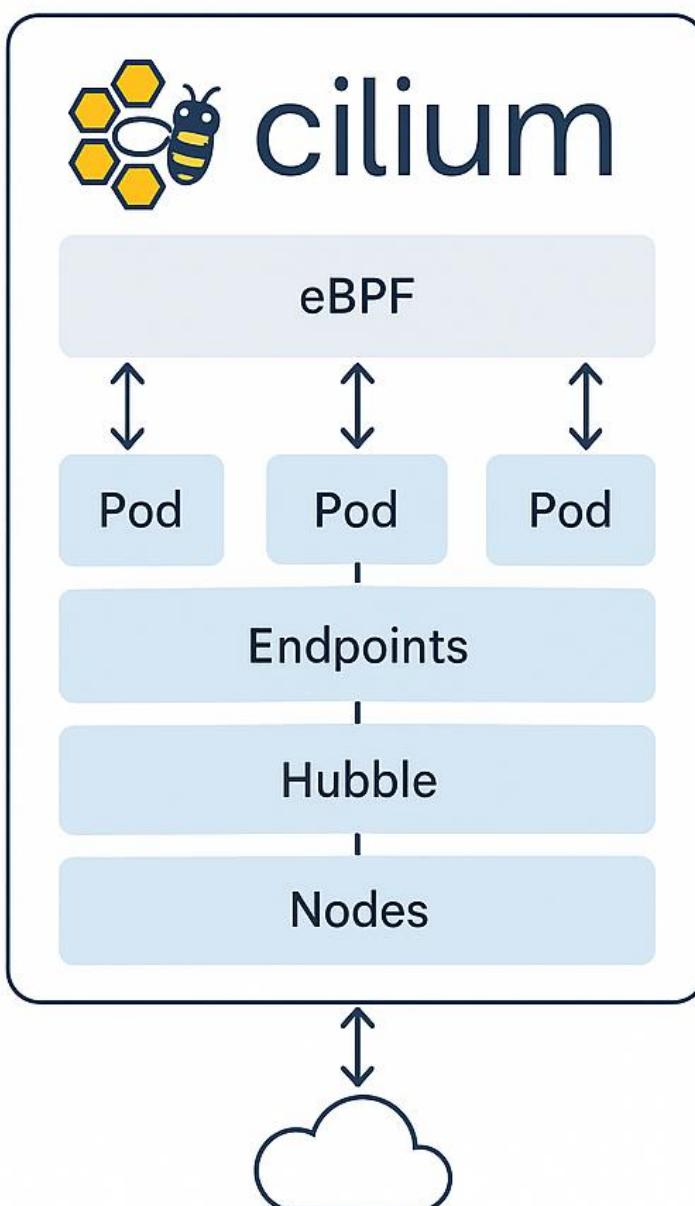


# Why eBPF Matters

| Traditional | eBPF                                    |
|-------------|---|
| Networking  | User-space → kernel transitions         |
| Performance | In-kernel processing                    |
| Use cases   | Load balancing, security, observability |
| Future      | The new networking standard             |

# Cilium: eBPF Networking Pioneer

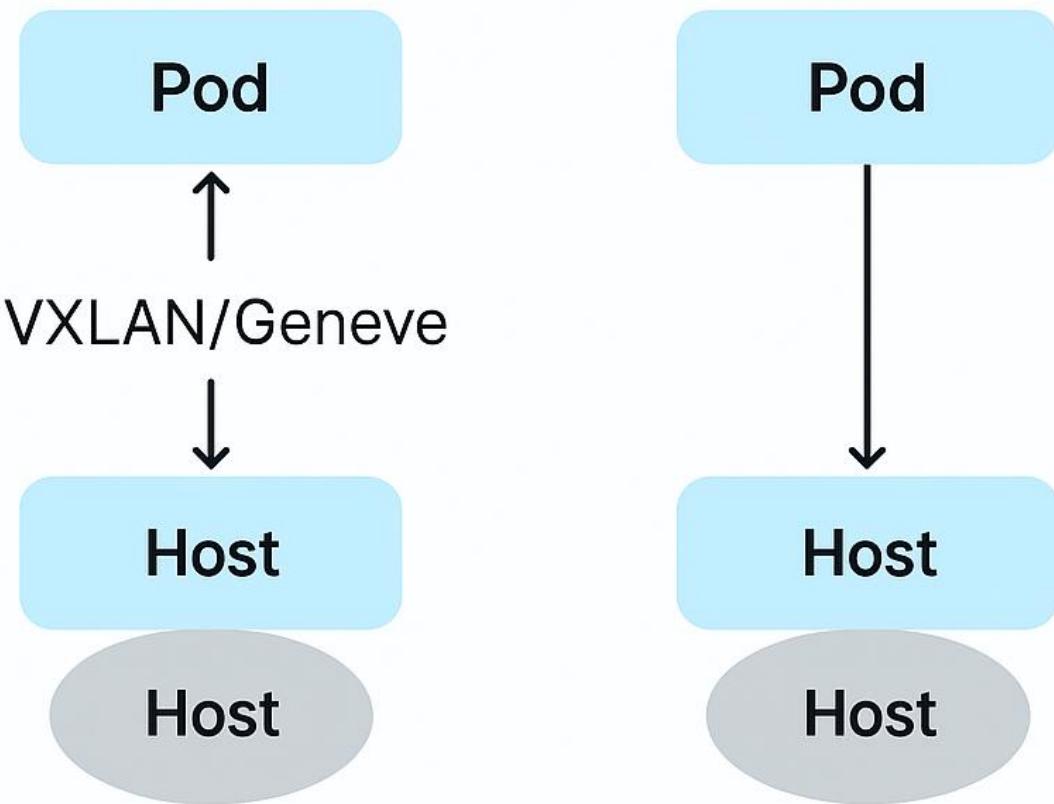
CNCF Project: Production-ready



- Identity-based security
- Load balancing without kube-proxy
- Advanced observability (Hubble)
- Multi-cluster networking

# Cilium Flexibility

|                            |                       |
|----------------------------|-----------------------|
| <b>Overlay Mode</b>        | <b>Native Routing</b> |
| VXLAN/Geneve encapsulation | Direct host routing   |



- Choice depends on:
  - Infrastructure capabilities
- Performance: Native > Overlay
- Complexity: Overlay < Native

# Load Balancing Reimagined

Traditional

kube-proxy

iptables

Cilium

eBPF  
hashtables

- Better performance
- Lower CPU usage
- Advanced algorithms (Maglev)

## Scale beyond limits

# Zero-Trust Networking

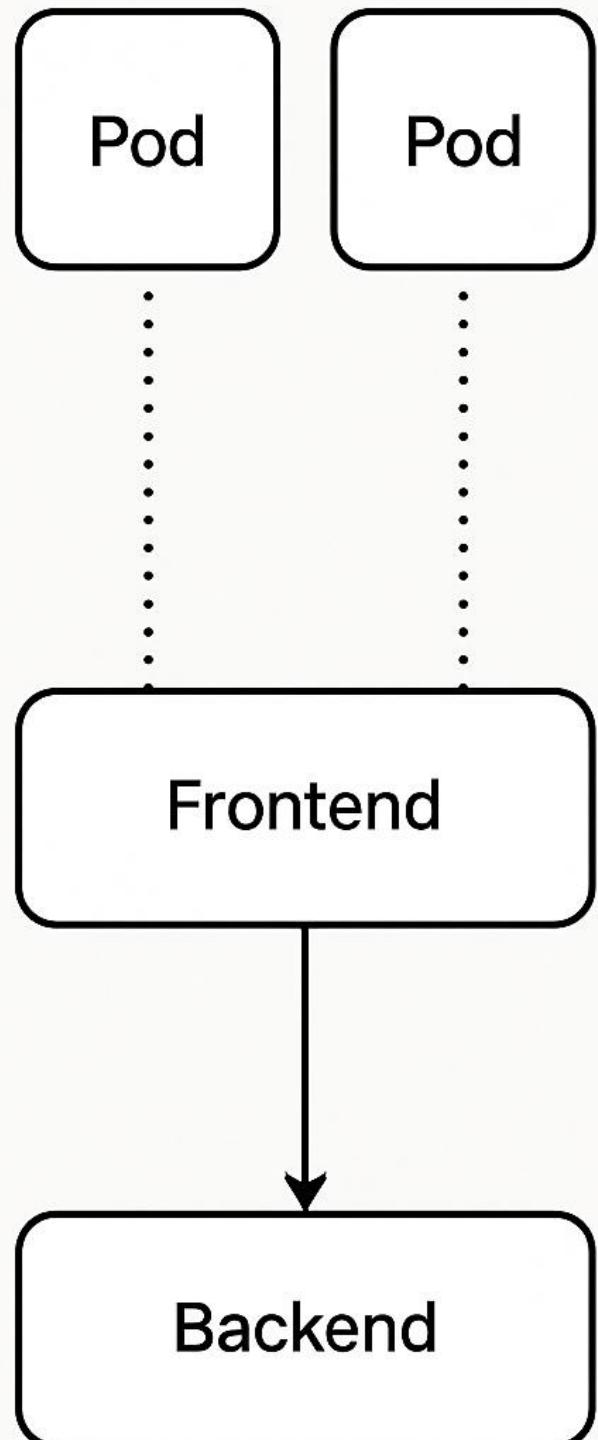
**Default:** All pods can communicate

**Network Policies:** Explicit allow rules

**Implementation:** Calico (reference)

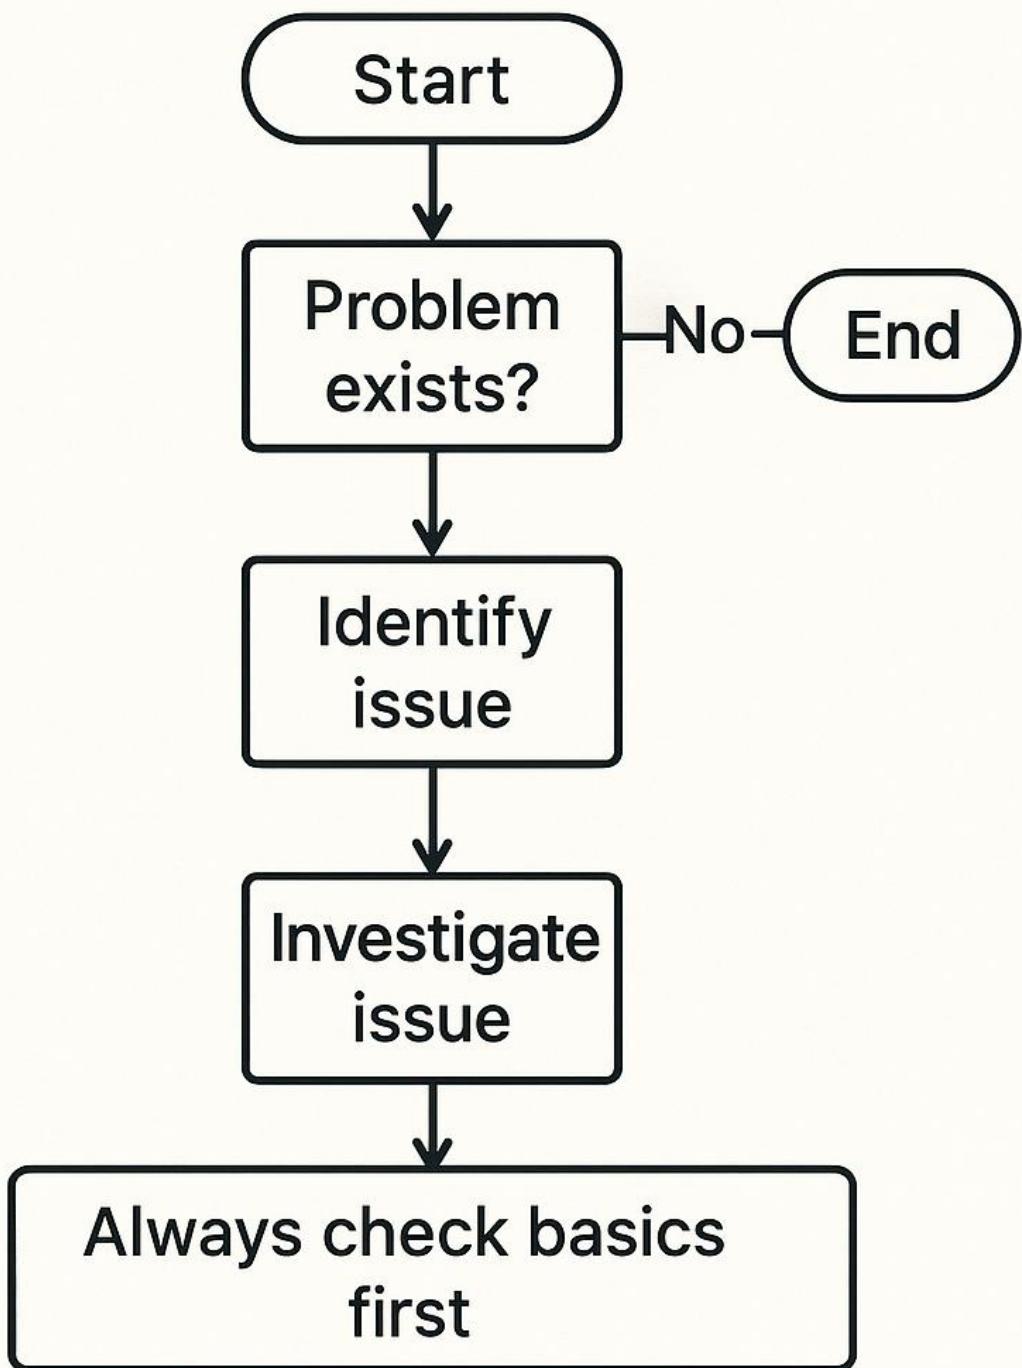
**Best Practice:** "Start with deny-all"

**Example:** Frontend → Backend only



# Network Debugging Toolkit

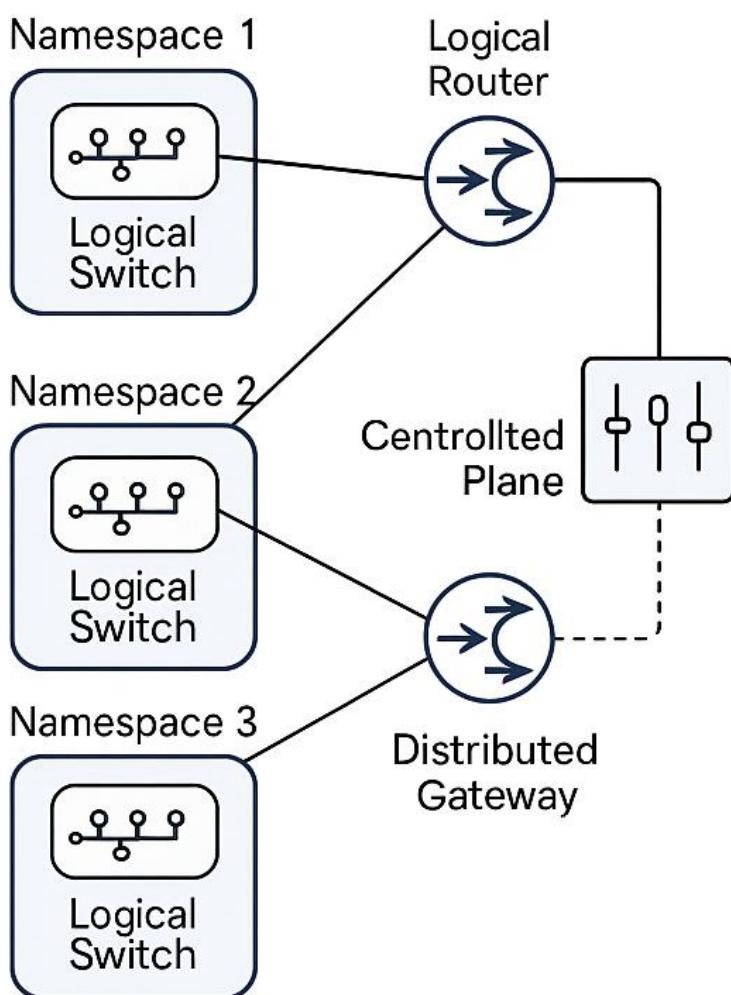
- kubectl exec for pod access
- dig/nslookup for DNS
- netstat/ss for connections
- tcpdump for packet capture



# OVN-Kubernetes: Software-Defined Networking

**What it is:** CNI plugin based on Open Virtual Network (OVN)

**Built on:** Open vSwitch (OVS) data plane



## Key Components

- Logical switches per namespace
- Logical routers for inter-namespace routing
- Distributed gateway for external access
- Network policies via CLs

## Architecture Benefits

- True multi-tenancy
- Advanced networking features
- Centralized control plane
- Hardware offload support

## Use Cases

Enterprise environments, multi-tenant clusters, complex networking requirements

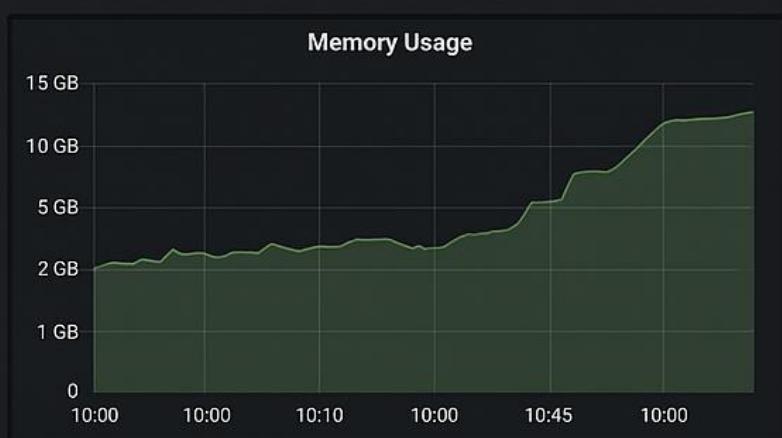
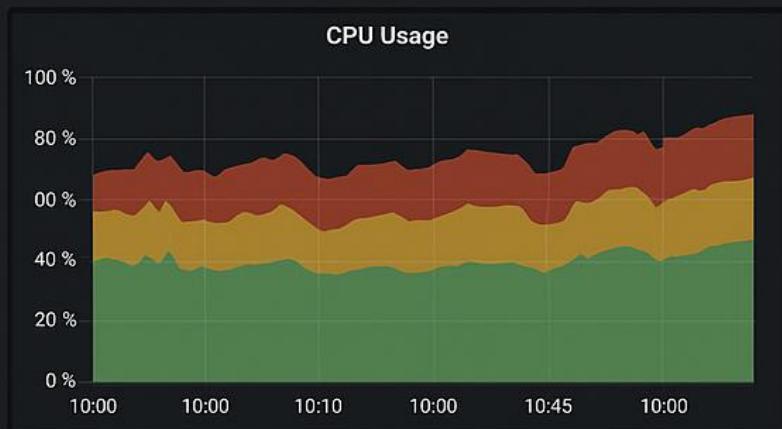
# OVN-K8s: Advanced Networking Features

| Enterprise Features   | Fix These Common Issues  |                      |
|---|--------------------------|----------------------|
| <ul style="list-style-type: none"><li>• Multiple external gateways</li><li>• Network segmentation</li><li>• Quality of Service (QoS)</li><li>• IPv6 dual-stack support</li><li>• Hardware acceleration</li><li>• DPDK integration</li></ul> | DNS resolution failures  | Check CoreDNS        |
| <h3>Network Policies</h3> <ul style="list-style-type: none"><li>• Default deny-all security</li><li>• Granular traffic control</li><li>• Stateful connection tracking</li></ul>   | Pod communication issues | Verify CNI           |
|   | Service unreachable      | Check endpoints      |
| <h3>Performance</h3> <ul style="list-style-type: none"><li>• Hardware offload capabilities</li><li>• DPDK for high throughput</li><li>• Smart NIC integration</li></ul>   | Slow performance         | Profile network path |
|   | MTU problems             | Check path MTU       |
| <b>Best for:</b> Telecom, financial services, large enterprises with complex security requirements  |                          |                      |

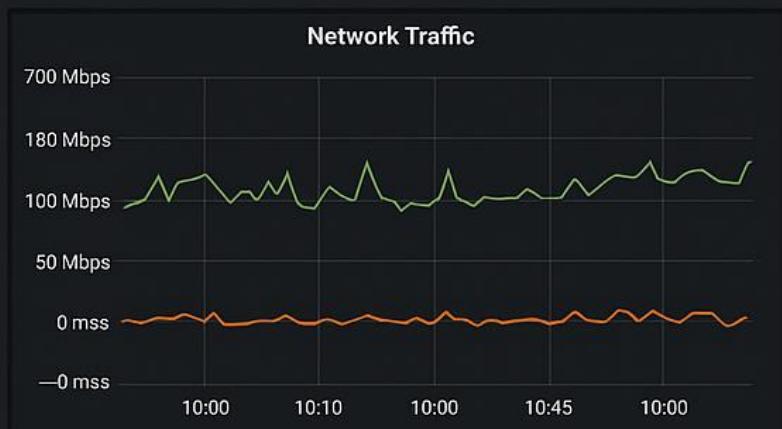
# Optimize for Performance

## Strategies:

- Choose right CNI plugin
- Tune MTU settings
- Use native routing when possible
- Consider eBPF solutions



## Monitoring: Prometheus + Grafana



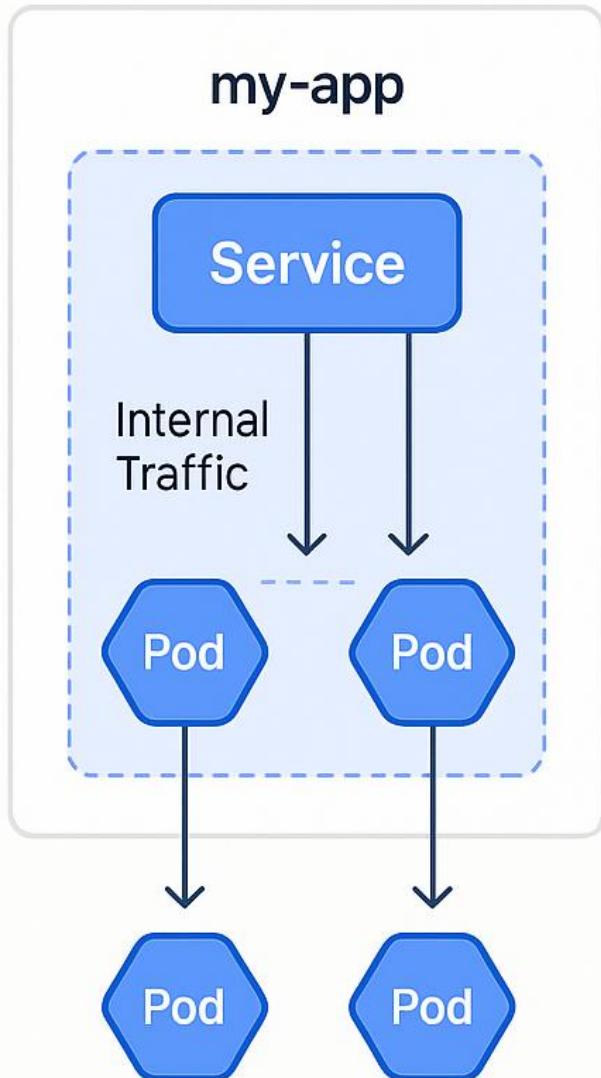
# Secure Your Network

-  Implement network policies
-  Use encrypted communication
-  Segment sensitive workloads
-  Monitor network traffic
-  Regular security audits

Falco for runtime security

# ClusterIP: Internal Communication Hub

Default service type for cluster-internal access



## Key Features:

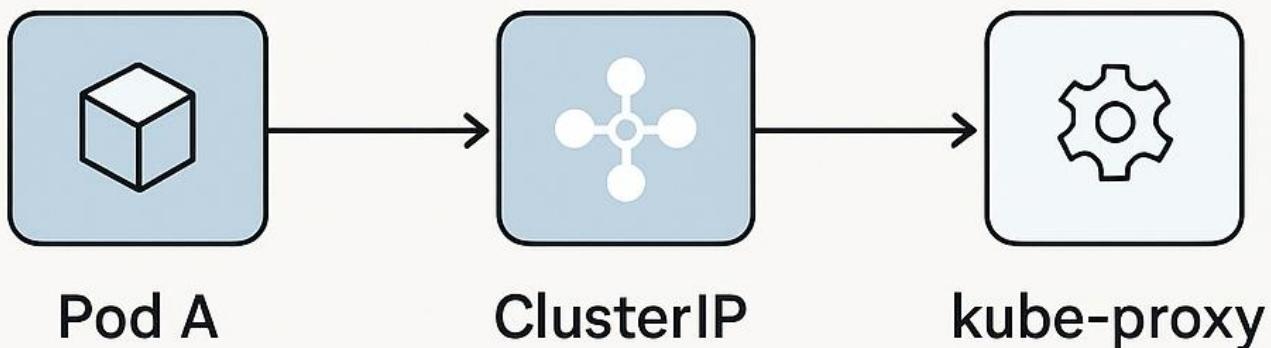
- Virtual IP accessible only within cluster
- Automatic load balancing to healthy pods
- DNS name: service-name.namespace svc.cluster.local
- Port mapping: Service port → Target port

## YAML Example:

```
apiVersion: v1
kind: Service
metadata:
  name: my-app
spec:
  type: ClusterIP
  selector:
    app: my-app
  ports
  - port: 80
    targetPort:8080
```

**Use Case:** Backend services, databases, internal APIs

# ClusterIP Traffic Flow



## Behind the scenes

- kube-proxy creates iptables rules
- Service IP is virtual (not assigned to any interface)
- Load balancing algorithms: Round-robin, least connections

## Troubleshooting

- Check endpoints: `kubectl get endpoints`
- Verify selector labels match
- Test from within cluster: `kubectl exec -it pod -- curl service-ip`

**Pro Tip:** Use service names, not IPs for better portability

# NodePort: External Access Gateway

Exposes service on each node's IP at static port

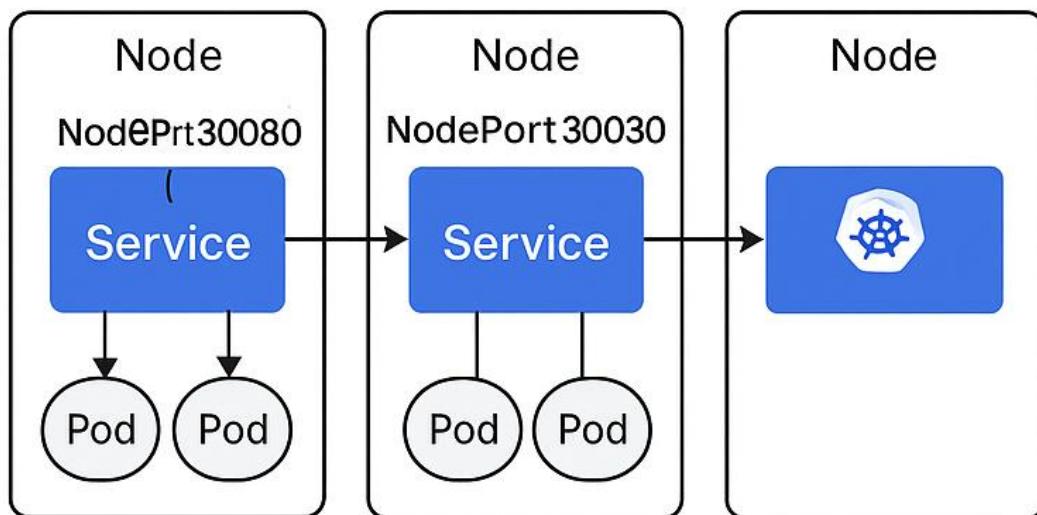
**Port Range:** 30000-32767 (configurable)

**Access Pattern:** NodeIP:NodePort → Service → Pocs

**Key Features:** External accessibility

High availability across nodes

Built on top of ClusterIP



- External accessibility
- High availability across nodes
- Built on top of ClusterIP

## YAML Example:

```
apiVersion: v1
kind: Service
metadata:
  name: my-web-api
spec:
  type: NodePort
  selector:
    app: web
  ports:
    - port: 80
      targetPort: 80
      nodePort: 3001
```

# NodePort Implementation Details

## How it works:

- Creates ClusterIP service
- Allocates port on ALL nodes
- Routes traffic from any node to healthy pods

## Traffic Flow



## Advantages

- Simple external access
- No external load balancer needed

## Limitations

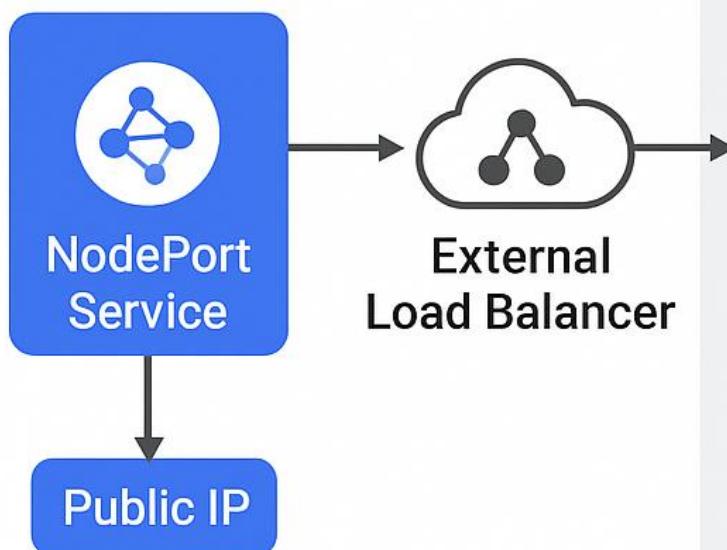
- Limited port range
- Manual port management
- No SSL termination

## Best for:

Development, small deployments,  
cost-conscious setups

# LoadBalancer: Cloud-Native External Access

- Cloud provider's load balancer integration
- Creates: NodePort service (autab)
- External load balancer
- Public IP assignment



## YAML Example

```
apiVersion: v1
kind: Service
metadata:
  name: web-service
spec:
  type: LoadBalancer
  selector:
    app: web
  ports:
    port: 80
    targetPort: 8080
```

**Cloud Support:** AWS ELB/ALB,  
GCP Load Balancer, Azure LoadBalancer

**Benefits:** Production-ready, scalable, managed

# ExternalName: DNS Alias Service

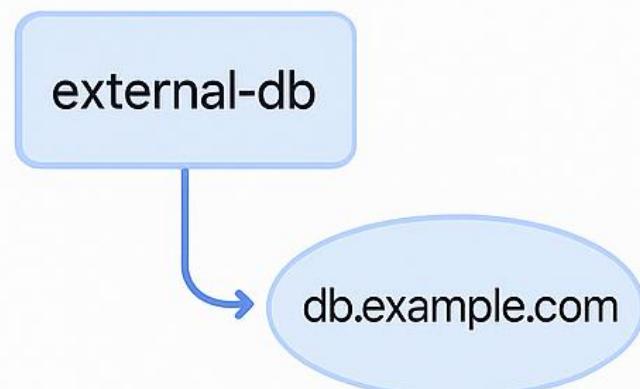
**Purpose:** Map service to external DNS name

No selectors or endpoints needed

**Use cases:**

- External database connections
- Third-party API integration
- Service migration scenarios

```
apiVersion: v1
kind: Service
metadata:
  name external-db
spec:
  type:
    ExternalName
  db.example.com
```



```
external-db.namespace.svc.cluster.local →
db.example.com
```

**Pro Tip:** Great for abstracting external dependencies

# Headless Services: Direct Pod Access

**Service with clusterIP: None**

## Behavior:

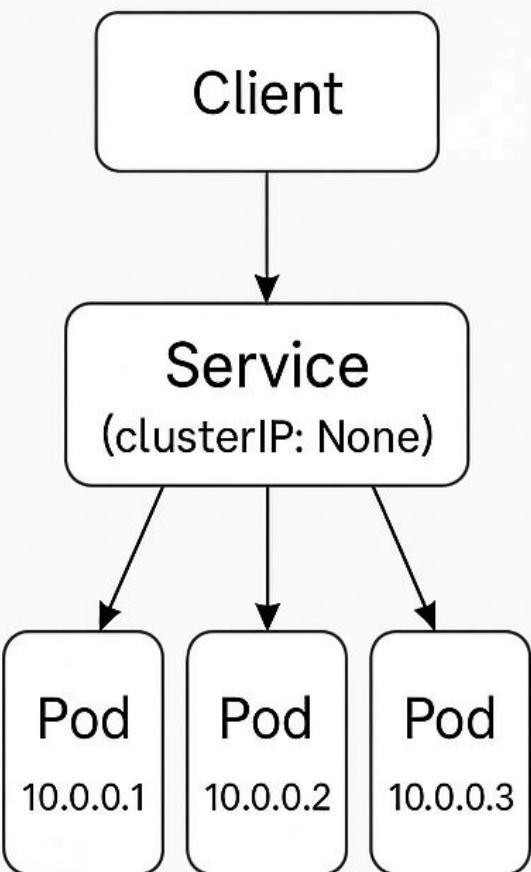
- No virtual IP assigned
- DNS returns pod IPs directly
- Client-side load balancing

## Use Cases:

- StatefulSets with stable network identities
- Custom load balancing logic
- Service discovery scenarios

## YAML Example:

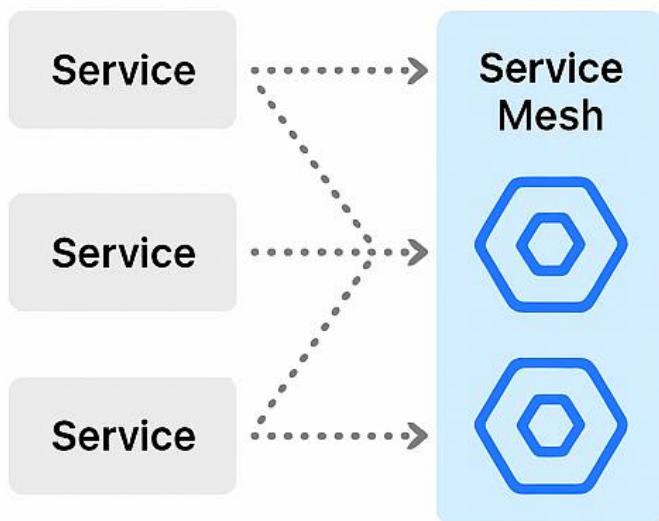
```
apiVersion: v1
kind: Service
metadata:
  name: headless-svc
spec:
  clusterIP: None
  selector:
    app: stateful-app
  ports
  - port: 80
```



## DNS Result:

Multiple A records  
for each pod

# Services + Service Mesh = 🚀



## What service mesh adds:

- Advanced traffic management
- Security (mTLS, policies)
- Observability and tracing
- Circuit breaking and retries

## Popular solutions:

- Istio (most feature-rich)
- Linkerd (lightweight)
- Consul Connect

## Benefits over basic services:

- Zero-trust security
- Traffic splitting
- Traffic splitting/canary deployments
- Distributed tracing

When to use: “Complex microservices architectures”

# Optimize Service Performance

## Key metrics to monitor

- Connection latency
- Throughput (request/sec)
- Error rates
- Pod health and readiness

## Throughput



## Connection Latency



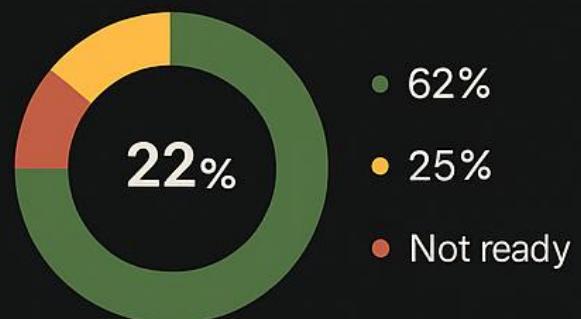
## Optimization strategies

- Tune session affinity
- Optimize health checks
- Use appropriate service
- Consider connection po-

## Tools

- Prometheus for metrics
- Jaeger for tracing
- Load testing with k6

## Pod Health and Readiness



**Goal:** Sub-100ms service response times

# Optimize Service Performance

## Key metrics to monitor

- Connection latency
- Throughput (requests/sec)
- Error rates
- Pod health and readiness

## Optimization strategies

- Tune session affinity
- Optimize health checks
- Use appropriate service type

## Tools

- Prometheus for metrics
- Jaeger for tracing
- Load testing with k6

Connection latency

51 ms



14:00 16:00 18:00 21:00

Throughput

1,365 requests/sec

Error rates

1.0%

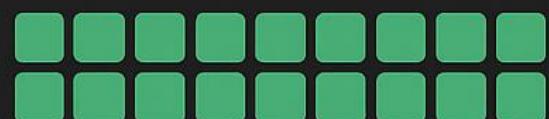
0.50

0:0%



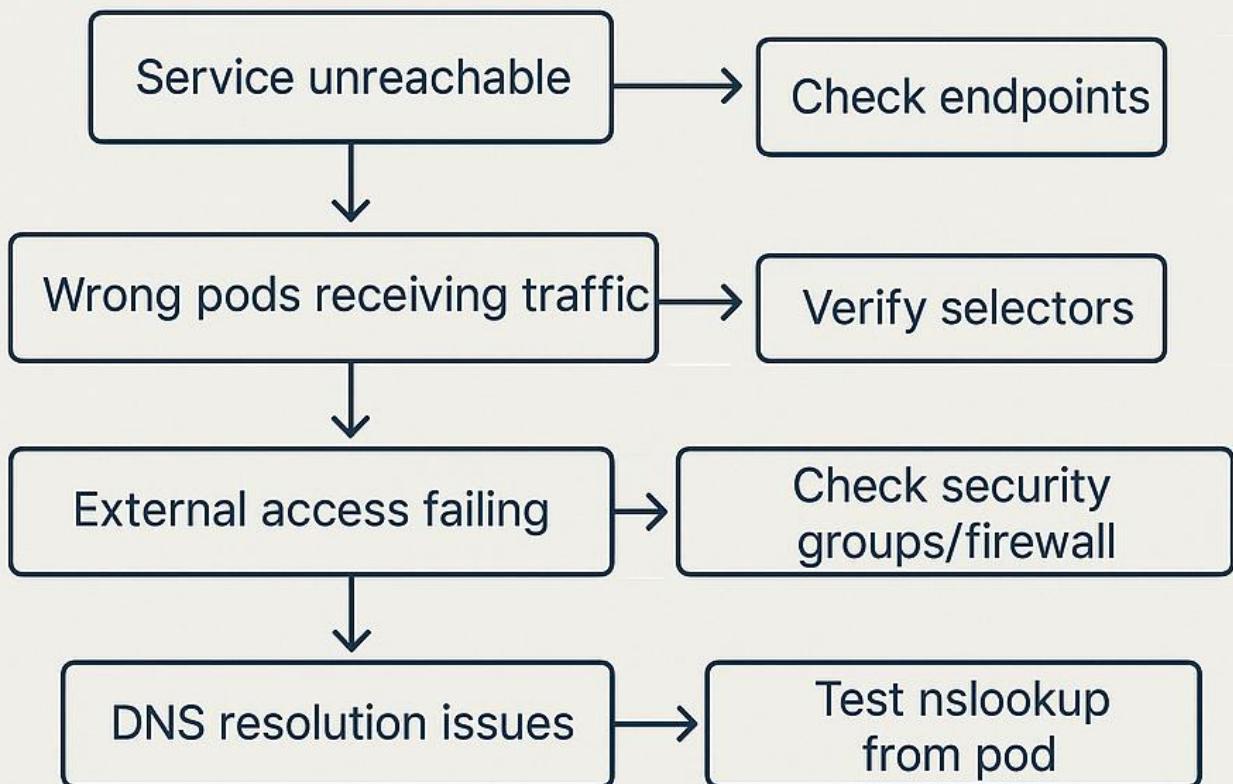
18:00 19:00 20:00 21:00 22:00

Pod health and readiness



## Sub-100ms service response times

# Debug Service Issues Like a Pro



## Debug commands:

```
bash kubectl get svc, endpoints  
kubectl describe service my-service  
kubectl logs -l app=my-app  
nslookup my-service.default.svc.cluster
```

## Pro tip:

Always verify pod labels match service selectors

# Production Readiness

- CNI plugin selected and tested
- Network policies implemented
- Service types chosen appropriately
- Load balancing strategy defined
- Monitoring and alerting setup
- Disaster recovery tested
- Performance baselines established
- Security scanning automated
- Service mesh evaluated (if needed)

# Ready to Go Deeper?



This was just the beginning of your K8s networking journey!



[Read my complete guide on Medium](https://medium.com/@salwan.mohamed/the-complete-kubernetes-networking-mastery-guide-from-basics-to-advanced-ebof-implementation-ebfddab65b2b)

<https://medium.com/@salwan.mohamed/the-complete-kubernetes-networking-mastery-guide-from-basics-to-advanced-ebof-implementation-ebfddab65b2b>\*

## What you'll get in the full guide:

- Step-by-step CNI plugin configurations
- Advanced eBPF implementation examples
- Production troubleshooting playbooks
- OVN-Kubernetes deployment guides
- Real-world case studies

**Join 10,000+ engineers mastering K8s networking**

👉 Like if this helped you

🔁 Repost to help your network

➕ Follow @salwan.mohamed for more K8s content

💬 Share your networking challenges below

Questions about your setup? Drop them in comments! ⤵

Save this post for your next K8s project 📖